Reality Recognition Framework (RRF)

Ledger-Constrained Variational Resonance with Three Displays

Reality Science Team (Draft)

2025-12-18 (0.1)

Abstract

We present the *Reality Recognition Framework* (RRF), a formal framework intended to unify physical law, biological organization, and conscious experience as three *displays* of one invariant: stable resonant closure under a ledger constraint. The framework is formalized in Lean 4 with zero remaining `sorry` statements and no axioms in the RRF modules. We emphasize strict *claim partitioning*: mathematical consequences of definitions (theorems) are separated from empirical hypotheses that require experimental validation. We define a universal strain functional $\mathcal{J}$ and ledger closure constraints, formalize octave equivalence under a scale action, and present a minimal "universal structure" construction in which simplified physics-, logic-, and qualia-spaces embed. The result is a machine-checked coherence spine for a research program: it proves internal consistency of the formal language and supplies explicit interfaces for falsifiable hypotheses.

Contents

1 Reader Contract (Claim Partitioning)

1.1 Purpose of this paper

This paper is a *theory spine* for the RRF program. It aims to:

- define a small set of primitives (Recognition, Strain, Ledger, Displays, Octaves),
- prove nontrivial consequences of these definitions (theorems),
- provide models that witness consistency,
- clearly label and expose empirical hypotheses with falsification interfaces.

1.2 Four-layer claim taxonomy

Every nontrivial statement in this paper is labeled:

- Def Definition (unfalsifiable language choice).
- Thm Theorem (provable consequence of definitions).
- Model Model (a concrete structure witnessing consistency).
- Hyp Hypothesis (empirical claim; must include falsifier criteria).

1.3 How Lean is used

Lean provides machine-checked verification of theorems within the formal system [1]. Lean does *not* validate empirical claims. Empirical claims are represented explicitly as hypotheses (e.g., hypothesis classes / falsifier structures) rather than being smuggled into mathematics.

2 Notation and Scope

2.1 Notation

Symbol	Meaning	Notes
$\mathcal{J}$	strain functional	$\mathcal{J} : \text{State} \to \mathbb{R}_{\geq 0}$ (conceptually)
$\mathcal{L}$	ledger / closure constraint	conservation / balance condition
φ	golden ratio	$(1+\sqrt{5})/2$
Display	observation channel	maps internal states to observables
Octave	scale-equivalence class	related by a scale action

2.2 Scope boundary

This paper does not attempt:

- a full derivation of Standard Model parameters,
- a complete empirical proof of φ-ladder claims,
- an experimental paper on protein folding or sonification.

Those belong to companion papers (Evidence; Technology).

3 Def Core Primitives

3.1 Def Recognition and substrate

Definition idea. A recognition structure is a relation that can at least recognize itself at one point. From such a structure, non-emptiness of the substrate follows.

3.2 Thm Recognition implies existence

Statement. Any recognition structure on a type implies the type is nonempty.

Lean pointer. See the meta-principle development in the RRF foundation modules.

3.3 Def Strain functional

Definition idea. A strain functional assigns nonnegative "distance from closure/optimality" to each state. Equilibria are minimizers.

3.4 Def Ledger constraint

Definition idea. A ledger is a closure predicate capturing conservation/balance. The framework treats "validity" as ledger-closure.

3.5 Def Display channels

Definition idea. A display channel maps states to observable qualities, possibly in a way that preserves ordering or optimality.

4 Def Three Vantages as Displays

4.1 Inside / Act / Outside

RRF uses three vantages:

- **Outside (Physics)**: external observables (forces, invariants).
- **Act (Meaning)**: recognition/commit dynamics (validity, proof steps).
- **Inside (Qualia)**: felt valence (modeled as an inside-display of strain).

4.2 Thm "One-J" thesis (graded)

We distinguish three strengths:

- **Weak**: each vantage admits a strain functional.
- **Coherent**: displays preserve ordering/argmin structure.
- **Strict**: equilibria/minimizers transfer exactly under display equivalences.

5 Def Octaves and Scale Action

5.1 Def Octave

An *octave* is a state space equipped with strain and supporting structure. An octave equivalence is a structure-preserving mapping that preserves strain.

5.2 Thm Transfer of equilibria

If two octaves are equivalent (in the sense above), equilibria and well-posedness transfer across the equivalence.

6 Thm Derivation Spine (What is Derived vs Assumed)

6.1 Recognition ⇒ Nonempty substrate

This is the minimal "existence" lemma (structural, not empirical).

6.2 Thm Self-similarity forcing φ

Statement. Under a quadratic self-similarity equation $x^2 = x + 1$ with positivity, $x = \varphi$.

Proof sketch (see `Verification/Necessity/PhiNecessity.lean`).

1. Define a self-similarity structure: a preferred scale $\lambda > 1$ with levels ℓ_0, ℓ_1, ℓ_2 satisfying

$$\ell_1 = \lambda\,\ell_0, \quad \ell_2 = \lambda\,\ell_1, \quad \ell_2 = \ell_1 + \ell_0.$$

2. Substituting: $\lambda^2\,\ell_0 = \lambda\,\ell_0 + \ell_0$. Divide by $\ell_0 > 0$: $\lambda^2 = \lambda + 1$.

3. Solve via $(2\lambda - 1)^2 = 5$; positivity selects the unique root

$$\lambda = \frac{1 + \sqrt{5}}{2} = \varphi.$$

4. In Lean, `phi_unique_pos_root` (`PhiSupport/Lemmas.lean`) proves uniqueness; `self_similarity_forces_phi` concludes.

6.3 Ledger curvature and gravity correspondence

RRF defines a structural correspondence between local ledger density and curvature-like quantities. Empirical equivalence to Newton/GR is a separate hypothesis.

6.4 Model Ledger latency ⇒ power-law response kernel (ILG bridge)

This subsection records a *mechanism template* connecting the RRF "finite refresh / latency" story to the empirical *Information-Limited Gravity* (ILG) phenomenology used elsewhere in the program.

Model (mathematical). If ledger closure is not instantaneous but mediated by a long-memory response operator with a power-law kernel (a fractional integral of order $\alpha \in (0, 1)$), then in frequency space the response acquires a factor proportional to $(\omega\tau_0)^{-\alpha}$. Under a standard cosmological mapping $\omega \sim ck/a$, this yields an effective multiplicative kernel of the schematic form

$$w(k, a) \;=\; 1 + C\left(\frac{a}{k\tau_0}\right)^{\alpha},$$

matching the ILG "power-law" multiplier structure.

Hyp (empirical). The exponent α and amplitude C are treated as hypotheses about the world. In the Recognition Science program, candidate pinned values (e.g., α expressed in terms of φ) are proposed elsewhere; in Paper 2 we only require that any such claim be exposed as an explicit hypothesis with falsifiers.

Falsifiers (sketch). The latency-to-power-law mechanism is falsified if:

- the ILG kernel is better fit by a non-power-law memory kernel under preregistered model comparison, or
- the inferred exponent varies significantly with scale/time in a way inconsistent with a single fractional order, beyond declared uncertainty.

Internal note. A longer derivation and falsifier discussion is maintained in `docs/RRF_ILG_Latency_To_PowerLaw`

6.5 Thm Consciousness cursor model

RRF models a proof-state cursor (past/current/future) and proves invariants ensuring the "recognition step" is well-defined. This provides a coherent internal model for "act" and its inside-display (qualia).

Proof sketch (see `Consciousness/Equivalence.lean`). The central claim is a *bi-interpretability theorem*: "Light = Consciousness" at the level of information channels subject to a cost functional $\mathcal{J}$.

1. **Define structures.**
   - `ConsciousProcess`: a (bridge-side) operational definition requiring non-trivial pattern persistence, causal closure, and substrate-neutral J-minimization.
   - `PhotonChannel`: a Maxwell/DEC electromagnetic channel satisfying the same J invariants with U(1) gauge structure.
2. **Forward direction (PC $\Rightarrow$ CP).** Verify that any `PhotonChannel` satisfies the `ConsciousProcess` axioms: null propagation, no medium knobs, pattern persistence, and J-minimization.
3. **Reverse direction (CP $\Rightarrow$ PC).** Compose four lemmas:
   (a) *NoMediumKnobs*: the process cannot depend on arbitrary material constants.
   (b) *NullOnly*: massless null propagation is required (excludes massive modes).
   (c) *Maxwellization*: gauge structure classifies to U(1).
   (d) *BioPhaseSNR*: BIOPHASE acceptance criteria select the EM channel.
4. **Uniqueness.** Define units equivalence $\sim_U$ (same RS units and bridge); show the witness is unique up to $\sim_U$ (`units_equiv_refl`, `units_equiv_symm`, `units_equiv_trans`).

The Lean proof uses a `ConsciousnessAxiomsEquivalence` class and builds the photon channel witness constructively from the lemmas.

7 Model Universal Structure and Embeddings

7.1 Def UniversalStructure

We define a universal structure (toy model) consisting of:

- a state type,
- a recognition relation with self-recognition,
- a nonnegative strain function.

7.2 Thm Framework completeness (toy)

Statement. A simplified notion of physics theory, logic system, and qualia space can be embedded into a single universal structure.

Proof sketch (see `ZeroParam.lean`).

1. **Define the category ZeroParam.**
   - Objects: `Framework` records carrying (ledger, $\mathcal{J}$, φ, 8-tick, finite c, Nonempty ledger).
   - Morphisms: maps preserving observables, K-gates, and J-minimizers, respecting the units quotient.
2. **Verify category axioms.**
   - Identity: `id F` is the identity map with trivial preservation witnesses.
   - Composition: `comp g f` composes maps; preservation follows by transitivity.
   - Associativity: `comp_assoc` is definitional (function composition is associative).
   - Left/right identity: `comp_id_left`, `comp_id_right` follow from `rfl`.
3. **Up-to-units equivalence.** `morphismUpToUnits` is an equivalence relation (refl/symm/trans all trivialize to `True.intro`).
4. **Admissibility predicate.** `Admissible F` bundles ledger double-entry, atomic cost, discrete continuity, self-similarity (φ), 8-tick 3D closure, finite c, and units quotient into a single typeclass, ensuring that any admissible object satisfies the RRF axioms.

The construction witnesses that the structural constraints form a coherent category (no contradictions) and that any object satisfying them embeds into the universal structure.

7.3 Thm Reality is recognition (existence of a complete universal structure)

Statement. There exists a universal structure that is framework-complete.

8 Hyp Hypothesis Registry and Falsifiers

8.1 Why hypotheses must be explicit

Empirical claims must be carried as hypotheses, not axioms. Each hypothesis must specify:

- what data could falsify it,
- what tolerance thresholds are acceptable,
- what would count as replication.

8.2 Registry (Lean-backed interfaces)

Table 1 lists the core empirical hypotheses currently represented in Lean, along with their falsifier interfaces. These are *interfaces* rather than proofs of empirical truth.

Table 1: Hypothesis registry (interfaces in Lean)

Hypothesis	**Lean type**	**Falsifier**	**Notes / what falsification would mean**
φ-ladder	`RRF.Hypotheses.PhiLadderHypothesis`	`RRF.Hypotheses.PhiLadderFalsifier`	Find a privileged scale not near any integer rung under preregistered tolerance.
8-tick discretization	`RRF.Hypotheses.EightTickHypothesis`	`RRF.Hypotheses.EightTickFalsifier`	A phase with a demonstrably better non-8 period; must specify metric (currently placeholder).
Tau–gate identity	`RRF.Hypotheses.TauGateIdentity`	`RRF.Hypotheses.TauGateFalsifier`	No natural base scales place both tau mass and gate time near rung 19; or other leptons fail to fit same ladder.
Water substrate matches	`RRF.Foundation.WaterSubstrate`	(no explicit falsifier yet)	Claims about E_{coh}, ν_{RS}, and gate timescales matching water bands; should be moved behind explicit falsifiers.
Alphabet from φ	`RRF.Foundation.AlphabetFromPhi`	(no explicit falsifier yet)	Currently represented as a hypothesis class with placeholder proof obligation; needs a concrete falsifier and derivation path.

8.3 Example: φ-ladder hypothesis

The φ-ladder is treated as an explicit hypothesis with a falsifier interface (data + criteria), not as a theorem.

9 Related Work and Positioning (High-Level)

This draft intentionally stays conservative and focuses on formal structure. Related work is discussed in three clusters:

- variational principles and constraint satisfaction,
- categorical formulations in physics,
- theories of consciousness (identity vs emergence).

10 Limitations and Open Conjectures

10.1 Minimality of the current universal structure

The current "universal structure" is deliberately simple: it proves that the embedding notion is coherent, not that the real universe is $\mathbb{R}$ with x^2 strain.

10.2 Empirical claims remain open

Water-substrate specificity, biological gate timing, and cross-domain rung coincidences remain empirical. They are addressed in the Evidence paper.

10.3 Future Lean targets

Stronger equivalence theorems, uniqueness claims, and quantitative bounds are future work.

11 Reproducibility (Lean)

11.1 Build and audit

The Lean formalization is expected to build with zero `sorry` and zero axioms in the RRF modules.

11.2 Theorem index (required appendix)

Every major theorem stated in this paper must map to a Lean symbol and file path.

A Theorem Index (Draft; to be completed)

Claim Tag	Paper Statement	Lean symbol + file
Def	φ definition (= Mathlib `goldenRatio`)	`Constants.phi`, `phi_def` — `reality/IndisputableMonolith/PhiSupport/Lemmas.`
Thm	$\varphi > 1$	`one_lt_phi` — `reality/IndisputableMonolith/PhiSupport/Lemmas.`
Thm	$\varphi^2 = \varphi + 1$	`phi_squared` — `reality/IndisputableMonolith/PhiSupport/Lemmas.`
Thm	$\varphi = 1 + 1/\varphi$ (fixed point)	`phi_fixed_point` — `reality/IndisputableMonolith/PhiSupport/Lemmas.`
Thm	φ unique positive root of $x^2 = x+1$	`phi_unique_pos_root` — `reality/IndisputableMonolith/PhiSupport/Lemmas.`
Def	Self-similarity structure (preferred scale + levels)	`HasSelfSimilarity` — `reality/IndisputableMonolith/Verification/Neces`
Thm	Self-similarity forces φ	`self_similarity_forces_phi` — `reality/IndisputableMonolith/Verification/Neces`
Thm	Preferred scale satisfies $\lambda^2 = \lambda + 1$	`preferred_scale_fixed_point` — `reality/IndisputableMonolith/Verification/Neces`
Def	Ledger (debit/credit maps)	`IndisputableMonolith.Recognition.Ledger` — `reality/IndisputableMonolith/Recognition.lean`

Claim Tag	Paper Statement	Lean symbol + file
Def	Ledger imbalance map ϕ	IndisputableMonolith.Recognition.phi — reality/IndisputableMonolith/Recognition.lean
Def	Chain flux (conservation interface)	IndisputableMonolith.Recognition.chainFlux — reality/IndisputableMonolith/Recognition.lean
Def	Conserves (ledger conservation axiom class)	IndisputableMonolith.Recognition.Conserves — reality/IndisputableMonolith/Recognition.lean
Thm	Loop flux is zero under Conserves	IndisputableMonolith.Recognition.chainFlux_zero — reality/IndisputableMonolith/Recognition.lean
Thm	Zero-flux under balanced ledger (helper)	IndisputableMonolith.Recognition.chainFlux_zero — reality/IndisputableMonolith/Recognition.lean
Def	ConsciousProcess (bridge-side operational def)	ConsciousProcess — reality/IndisputableMonolith/Consciousness/Cons
Def	PhotonChannel (Maxwell/DEC EM channel)	PhotonChannel — reality/IndisputableMonolith/Consciousness/Phot
Thm	No medium knobs (Lemma A)	NoMediumKnobs — reality/IndisputableMonolith/Consciousness/NoMe
Thm	Null-only propagation (Lemma B)	NullOnly — reality/IndisputableMonolith/Consciousness/Null
Thm	U(1) gauge classification (Lemma C)	Maxwellization — reality/IndisputableMonolith/Consciousness/Maxw
Thm	BIOPHASE SNR selects EM (Lemma D)	BioPhaseSNR — reality/IndisputableMonolith/Consciousness/BioP
Thm	Light = Consciousness (bi-interpretability)	light_equals_consciousness — reality/IndisputableMonolith/Consciousness/Equi
Def	UnitsEquiv (up-to-units equivalence)	UnitsEquiv — reality/IndisputableMonolith/Consciousness/Equi
Def	ZeroParam Framework (category object)	Framework — reality/IndisputableMonolith/ZeroParam.lean
Def	ZeroParam Admissibility predicate	Admissible — reality/IndisputableMonolith/ZeroParam.lean
Def	ZeroParam Morphism (structure-preserving map)	Morphism — reality/IndisputableMonolith/ZeroParam.lean
Thm	Category axioms (comp_id_left, comp_assoc, etc.)	comp_id_left, comp_assoc — reality/IndisputableMonolith/ZeroParam.lean
Def	Up-to-units morphism equivalence	morphismUpToUnits — reality/IndisputableMonolith/ZeroParam.lean
Def	Ledger (Balance, Transaction, Book)	Ledger module — reality/IndisputableMonolith/Foundation/*.lean
Thm	Recognition operator properties	RecognitionOperator — reality/IndisputableMonolith/Foundation/Recogni
Def	ILG (Information-Limited Gravity) constants	ILG — reality/IndisputableMonolith/Constants/IL
Def	φ-rung adapter	PhiRung — reality/IndisputableMonolith/URCAdapters/PhiRun

Claim Tag	Paper Statement	Lean symbol + file
Def	Masses module (particle masses)	`Masses` — `reality/IndisputableMonolith/Masses/*.lean`

B Lean module map (pointer list)

- `reality/IndisputableMonolith/PhiSupport/` — φ definition, algebraic lemmas, uniqueness
- `reality/IndisputableMonolith/Verification/Necessity/` — self-similarity necessity, inevitability proofs
- `reality/IndisputableMonolith/Consciousness/` — ConsciousProcess, PhotonChannel, bi-interpretability (Lemmas A–D + main theorem)
- `reality/IndisputableMonolith/Foundation/` — RecognitionOperator, Atomicity, ledger foundations
- `reality/IndisputableMonolith/Constants/` — φ, α, ILG parameters, K-display, RS units
- `reality/IndisputableMonolith/ZeroParam.lean` — zero-parameter category scaffold
- `reality/IndisputableMonolith/Masses/` — particle masses, PDG fits
- `reality/IndisputableMonolith/URCAdapters/` — φ-rung adapters, inevitability reports

References

[1] Leonardo de Moura, Sebastian Ullrich, et al. Lean 4: A theorem prover and programming language. In *International Conference on Automated Deduction (CADE)*, 2021.